


Java Threads

Multitasking and Multithreading

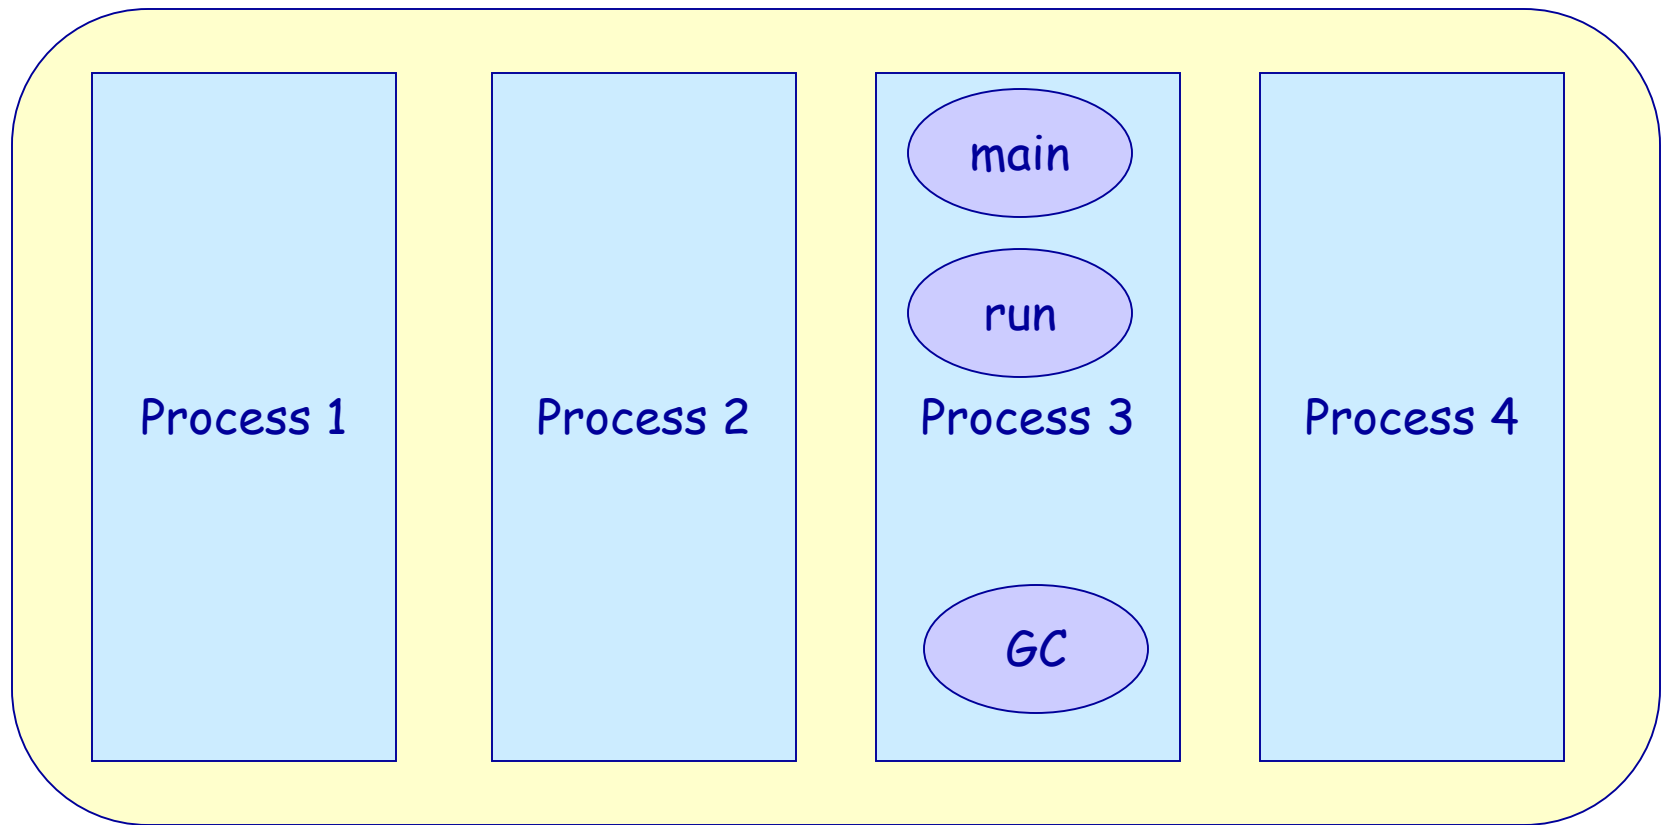
- **Multitasking:**
 - refers to a computer's ability to perform multiple jobs concurrently
 - more than one program are running concurrently, e.g., UNIX
- **Multithreading:**
 - A **thread** is a single sequence of execution within a program
 - refers to multiple threads of control within a single program
 - each program can run multiple threads of control within it, e.g., Web Browser

Concurrency vs. Parallelism



Threads and Processes

CPU



What are Threads Good For?

- To maintain responsiveness of an application during a long running task
- To enable cancellation of separable tasks
- Some problems are intrinsically parallel
- To monitor status of some resource (e.g., DB)
- Some APIs and systems demand it (e.g., Swing)

Application Thread

- When we execute an application:
 1. The JVM **creates** a Thread object whose task is defined by the `main()` method
 2. The JVM **starts** the thread
 3. The thread **executes** the statements of the program one by one
 4. After executing all the statements, the method returns and the **thread dies**

Multiple Threads in an Application

- Each thread has its private run-time stack
- If two threads execute the same method, each will have its own copy of the local variables the methods uses
- However, all threads see the same dynamic memory, i.e., heap (are there variables on the heap?)
- Two different threads can act on the same object and same static fields concurrently

Creating Threads

- There are two ways to create our own Thread object
 1. Subclassing the **Thread** class and instantiating a new object of that class
 2. Implementing the **Runnable** interface
- In both cases the `run()` method should be implemented

Extending Thread

```
public class ThreadExample extends Thread {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println("---");  
        }  
    }  
}
```

Thread Methods

void start()

- Creates a new thread and makes it runnable
- This method can be called only once

void run()

- The new thread begins its life inside this method

void stop() (deprecated)

- The thread is being terminated

Thread Methods

void yield()

- Causes the currently executing thread object to temporarily pause and allow other threads to execute
- Allow only threads of the same priority to run

void sleep(int *m*) or sleep(int *m*, int *n*)

- The thread sleeps for *m* milliseconds, plus *n* nanoseconds

Implementing Runnable

```
public class RunnableExample implements Runnable {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println ("***");  
        }  
    }  
}
```

A Runnable Object

- When running the Runnable object, a Thread object is created from the Runnable object
- The Thread object's run() method calls the Runnable object's run() method
- Allows threads to run inside any object, regardless of inheritance

Example - an applet
that is also a thread


Starting the Threads

```
public class ThreadsStartExample {  
    public static void main (String argv[]) {  
        new ThreadExample ().start ();  
        new Thread(new RunnableExample ().start ());  
    }  
}
```

What will we see when running
ThreadsStartExample?

```
mangal: /cs/phd/yarok/java/dbi/threads
[yarok@mangal]~/java/dbi/threads> java ThreadsStartExample
-----***-----
-----***-----***
-----***-----*****
*****
*****
*****
*****
[yarok@mangal]~/java/dbi/threads>
```

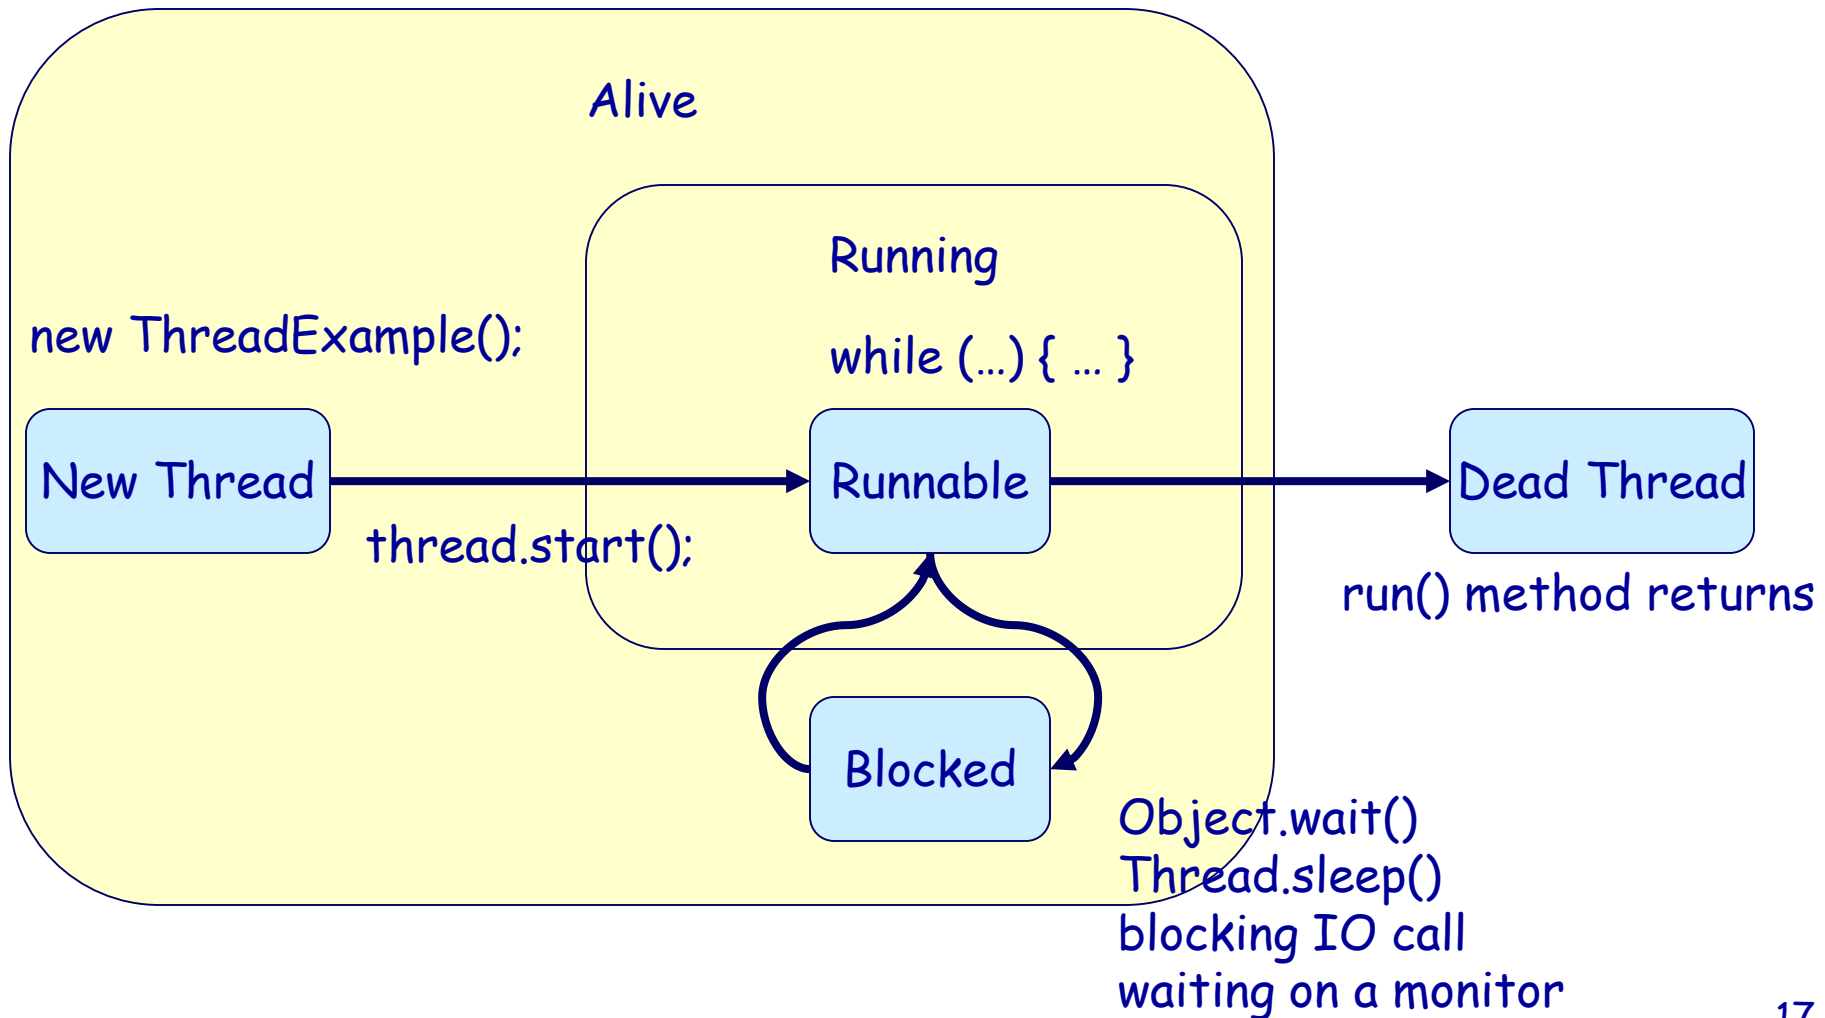
Scheduling Threads



What happens when a program with a `ServerSocket` calls `accept()`?

- Waiting for I/O operation to be completed
- Waiting to be notified
- Sleeping
- Waiting to enter a synchronized section

Thread State Diagram



Example

```
public class PrintThread1 extends Thread {  
    String name;  
  
    public PrintThread1(String name) {  
        this.name = name;  
    }  
  
    public void run() {  
        for (int i=1; i<100 ; i++) {  
            try {  
                sleep((long) (Math.random() * 100)) ;  
            } catch (InterruptedException ie) { }  
            System.out.print(name) ;  
        }  
    }  
}
```

Example (cont)

```
public static void main(String args[]) {
```

```
    PrintThread1 a = new PrintThread1("*");
```

```
    PrintThread1 b = new PrintThread1("-");
```

```
    a.start();
```

```
    b.start();
```

```
}
```

```
}
```

```
mangal:/cs/phd/yarok/java/dbi/threads
[yarok@mangal]~/java/dbi/threads> java ThreadsTest1
***-----***
--***--***--***--***--***--***--***--***--***--***
--***--***--***--***--***--***--***--***--***--***
--***--***--***--***--***--***--***--***--***--***
***--***--***--***--***--***--***--***--***--***
[yarok@mangal]~/java/dbi/threads>
```

Scheduling

- Thread **scheduling** is the mechanism used to determine how runnable threads are allocated CPU time
- A thread-scheduling mechanism is either **preemptive** or **nonpreemptive**

Preemptive Scheduling

- **Preemptive scheduling** - the thread scheduler preempts (pauses) a running thread to allow different threads to execute
- **Nonpreemptive scheduling** - the scheduler never interrupts a running thread
- The **nonpreemptive scheduler** relies on the running thread to yield control of the CPU so that other threads may execute

Starvation

- A nonpreemptive scheduler may cause **starvation** (runnable threads, ready to be executed, wait to be executed in the CPU a very long time, maybe even forever)
- Sometimes, starvation is also called a **livelock**

Time-Sliced Scheduling

- Time-sliced scheduling
 - the scheduler allocates a period of time that each thread can use the CPU
 - when that amount of time has elapsed, the scheduler preempts the thread and switches to a different thread
- Nontime-sliced scheduler
 - the scheduler does not use elapsed time to determine when to preempt a thread
 - it uses other criteria such as priority or I/O status

Java Scheduling

- Scheduler is preemptive and based on priority of threads
- Uses fixed-priority scheduling:
 - Threads are scheduled according to their priority w.r.t. other threads in the ready queue

Java Scheduling

- The highest priority runnable thread is always selected for execution above lower priority threads
- When multiple threads have equally high priorities, only one of those threads is guaranteed to be executing
- Java threads are guaranteed to be preemptive-but not time sliced
- **Q:** Why can't we guarantee time-sliced scheduling?

What is the danger of such scheduler?

Thread Priority

- Every thread has a priority
- When a thread is created, it inherits the priority of the thread that created it
- The priority values range from 1 to 10, in increasing priority

Thread Priority (cont.)

- The priority can be adjusted subsequently using the `setPriority()` method
- The priority of a thread may be obtained using `getPriority()`
- Priority constants are defined:
 - `MIN_PRIORITY=1`
 - `MAX_PRIORITY=10`
 - `NORM_PRIORITY=5`

The **main** thread is created with priority `NORM_PRIORITY`

Some Notes

- Thread implementation in Java is actually based on operating system support
- Some Windows operating systems support only 7 priority levels, so different levels in Java may actually be mapped to the same operating system level
- What should we do about this?
- Furthermore, The thread scheduler may choose to run a lower priority thread to avoid starvation

Daemon Threads

- **Daemon** threads are “background” threads, that provide services to other threads, e.g., the garbage collection thread
- The Java VM **will not exit** if non-Daemon threads are executing
- The Java VM **will exit** if only Daemon threads are executing
- Daemon threads die when the Java VM exits
- **Q:** Is the **main** thread a daemon thread?

Thread and the Garbage Collector

- Can a Thread object be collected by the garbage collector while running?
 - If not, why?
 - If yes, what happens to the execution thread?
- When can a Thread object be collected?


ThreadGroup

- The ThreadGroup class is used to create groups of similar threads. Why is this needed?

“Thread groups are best viewed as an unsuccessful experiment, and you may simply ignore their existence.”

Joshua Bloch, software architect at Sun

Multithreading Client-Server



Server

```
import java.net.*;import java.io.*;
class HelloServer {

    public static void main(String[] args) {
        int port = Integer.parseInt(args[0]);
        try {
            ServerSocket server =
                new ServerSocket(port);
        } catch (IOException ioe) {
            System.err.println("Couldn't run " +
                "server on port " + port);
            return;
        }
    }
}
```

```
while(true) {  
    try {  
        Socket connection = server.accept();  
        ConnectionHandler handler =  
            new ConnectionHandler(connection);  
        new Thread(handler).start();  
    } catch (IOException ioel) {  
    }  
}
```

Connection Handler

```
// Handles a connection of a client to an  
HelloServer.  
// Talks with the client in the 'hello' protocol  
class ConnectionHandler implements Runnable {  
  
    // The connection with the client  
    private Socket connection;  
  
    public ConnectionHandler(Socket connection) {  
        this.connection = connection;  
    }  
}
```

```
public void run() {  
    try {  
        BufferedReader reader =  
            new BufferedReader(  
                new InputStreamReader(  
                    connection.getInputStream()) );  
  
        PrintWriter writer =  
            new PrintWriter(  
                new OutputStreamWriter(  
                    connection.getOutputStream()) );  
  
        String clientName = reader.readLine();  
        writer.println("Hello " + clientName);  
        writer.flush();  
    } catch (IOException ioe) {}  
}  
}
```

Client side

```
import java.net.*; import java.io.*;
```

```
// A client of an HelloServer
```

```
class HelloClient {
```

```
    public static void main(String[] args) {
```

```
        String hostname = args[0];
```

```
        int port = Integer.parseInt(args[1]);
```

```
        Socket connection = null;
```

```
        try {
```

```
            connection = new Socket(hostname, port);
```

```
        } catch (IOException ioe) {
```

```
            System.err.println("Connection failed");
```

```
            return;
```

```
        }
```

```
try {  
    BufferedReader reader =  
        new BufferedReader(  
            new InputStreamReader(  
                connection.getInputStream() ) ) ;  
    PrintWriter writer =  
        new PrintWriter(  
            new OutputStreamWriter(  
                connection.getOutputStream() ) ) ;  
  
    writer.println(args[2]); // client name  
    String reply = reader.readLine();  
    System.out.println("Server reply: "+reply);  
    writer.flush();  
} catch (IOException ioel) {  
}  
}
```

Note that the Client has not
changed from the
networking-lecture example 40


Concurrency

- An object in a program can be changed by more than one thread
- **Q:** Is the order of changes that were performed on the object important?

Race Condition

- A **race condition** - the outcome of a program is affected by the order in which the program's threads are allocated CPU time
- Two threads are simultaneously modifying a single object
- Both threads “race” to store their value

Race Condition Example



Monitors

- Each object has a “**monitor**” that is a token used to determine which application thread has control of a particular object instance
- In execution of a **synchronized** method (or block), access to the object monitor must be gained before the execution
- Access to the object monitor is queued

Monitor (cont.)

- Entering a monitor is also referred to as **locking** the monitor, or **acquiring ownership** of the monitor
- If a thread A tries to acquire ownership of a monitor and a different thread has already entered the monitor, the current thread (A) must wait until the other thread leaves the monitor

Critical Section

- The synchronized methods define **critical sections**
- Execution of critical sections is mutually exclusive. Why?

Example

```
public class BankAccount {
```


```
    private float balance;
```

```
    public synchronized void deposit(float amount) {  
        balance += amount;  
    }
```

```
    public synchronized void withdraw(float amount) {  
        balance -= amount;  
    }
```

```
}
```

Critical Sections



Java Locks are Reentrant

- Is there a problem with the following code?

```
public class Test {  
    public synchronized void a() {  
        b();  
        System.out.println("I am at a");  
    }  
    public synchronized void b() {  
        System.out.println("I am at b");  
    }  
}
```

Static Synchronized Methods

- Marking a static method as synchronized, associates a monitor with the class itself
- The execution of synchronized static methods of the same class is mutually exclusive. Why?

Synchronized Statements

- A monitor can be assigned to a block:

```
synchronized(object) { some-code }
```

- It can also be used to monitor access to a data element that is not an object, e.g., array:

```
void arrayShift(byte[] array, int count) {  
    synchronized(array) {  
        System.arraycopy (array, count, array,  
                           0, array.size - count);  
    }  
}
```

The Followings are Equivalent

```
public synchronized void a() {  
  
    //... some code ...  
  
}
```

```
public void a() {  
  
    synchronized (this) {  
  
        //... some code ...  
  
    }  
  
}
```

The Followings are Equivalent

```
public static synchronized void a() {  
    //... some code ...  
}
```

```
public void a() {  
    synchronized (this.getClass()) {  
        //... some code ...  
    }  
}
```

Example

```
public class MyPrinter {
```

```
    public MyPrinter() {}
```

```
    public synchronized void printName(String name) {
```

```
        for (int i=1; i<100 ; i++) {
```

```
            try {
```

```
                Thread.sleep((long) (Math.random() * 100));
```

```
            } catch (InterruptedException ie) {}
```

```
        System.out.print(name);
```

```
    }
```

```
}
```

```
}
```

Example

```
public class PrintThread2 extends Thread {
```

```
    String name;
```

```
    MyPrinter printer;
```

```
    public PrintThread2(String name, MyPrinter printer){
```

```
        this.name = name;
```

```
        this.printer = printer;
```

```
    }
```

```
    public void run() {
```

```
        printer.printName(name) ;
```

```
    }
```

```
}
```

Example (cont)

```
public class ThreadsTest2 {  
  
    public static void main(String args[]) {
```

```
        MyPrinter myPrinter = new MyPrinter();
```

```
        PrintThread2 a = new PrintThread2("*", printer);
```

```
        PrintThread2 b = new PrintThread2("-", printer);
```

```
        PrintThread2 c = new PrintThread2("=", printer);
```

```
        a.start();
```

```
        b.start();
```

```
        c.start();
```

```
    }
```

```
}
```

What will happen?


```
mangal:/cs/phd/yarok/java/dbi/threads
[ yarok@mangal ] ~/java/dbi/threads> java ThreadsTest2

*****
*****-----
-----
-----
=====
=====
=====
=====
[ yarok@mangal ] ~/java/dbi/threads>
```

Deadlock Example

```
public class BankAccount {
```

```
    private float balance;
```

```
    public synchronized void deposit(float amount) {  
        balance += amount;  
    }
```

```
    public synchronized void withdraw(float amount) {  
        balance -= amount;  
    }
```

```
    public synchronized void transfer  
        (float amount, BankAccount target) {  
        withdraw(amount);  
        target.deposit(amount);  
    }
```

```
}
```

```
public class MoneyTransfer implements Runnable {
```

```
    private BankAccount from, to;  
    private float amount;
```

```
    public MoneyTransfer(  
        BankAccount from, BankAccount to, float amount) {  
        this.from = from;  
        this.to = to;  
        this.amount = amount;  
    }
```

```
    public void run() {  
        source.transfer(amount, target);  
    }
```


```
}
```

```
BankAccount aliceAccount = new BankAccount();  
BankAccount bobAccount = new BankAccount();  
...
```

```
// At one place  
Runnable transaction1 =  
    new MoneyTransfer(aliceAccount, bobAccount, 1200);  
Thread t1 = new Thread(transaction1);  
t1.start();
```

```
// At another place  
Runnable transaction2 =  
    new MoneyTransfer(bobAccount, aliceAccount, 700);  
Thread t2 = new Thread(transaction2);  
t2.start();
```

Deadlocks



Thread Synchronization

- We need to synchronized between transactions, for example, the consumer-producer scenario



Wait and Notify

- Allows two threads to cooperate
- Based on a single shared lock object
 - Marge put a cookie wait and notify Homer
 - Homer eat a cookie wait and notify Marge
 - Marge put a cookie wait and notify Homer
 - Homer eat a cookie wait and notify Marge

The wait() Method

- The **wait()** method is part of the `java.lang.Object` interface
- It requires a lock on the object's monitor to execute
- It must be called from a synchronized method, or from a synchronized segment of code. Why?

The wait() Method

- wait() causes the current thread to wait until another thread invokes the **notify()** method or the **notifyAll()** method for this object
- Upon call for wait(), the thread releases ownership of this monitor and waits until another thread notifies the waiting threads of the object

The **wait()** Method

- **wait()** is also similar to **yield()**
 - Both take the current thread off the execution stack and force it to be rescheduled
- However, **wait()** is not automatically put back into the scheduler queue
 - **notify()** must be called in order to get a thread back into the scheduler's queue
 - The objects monitor must be reacquired before the thread's run can continue

What is the difference between **wait** and **sleep**?

Consumer

- Consumer:

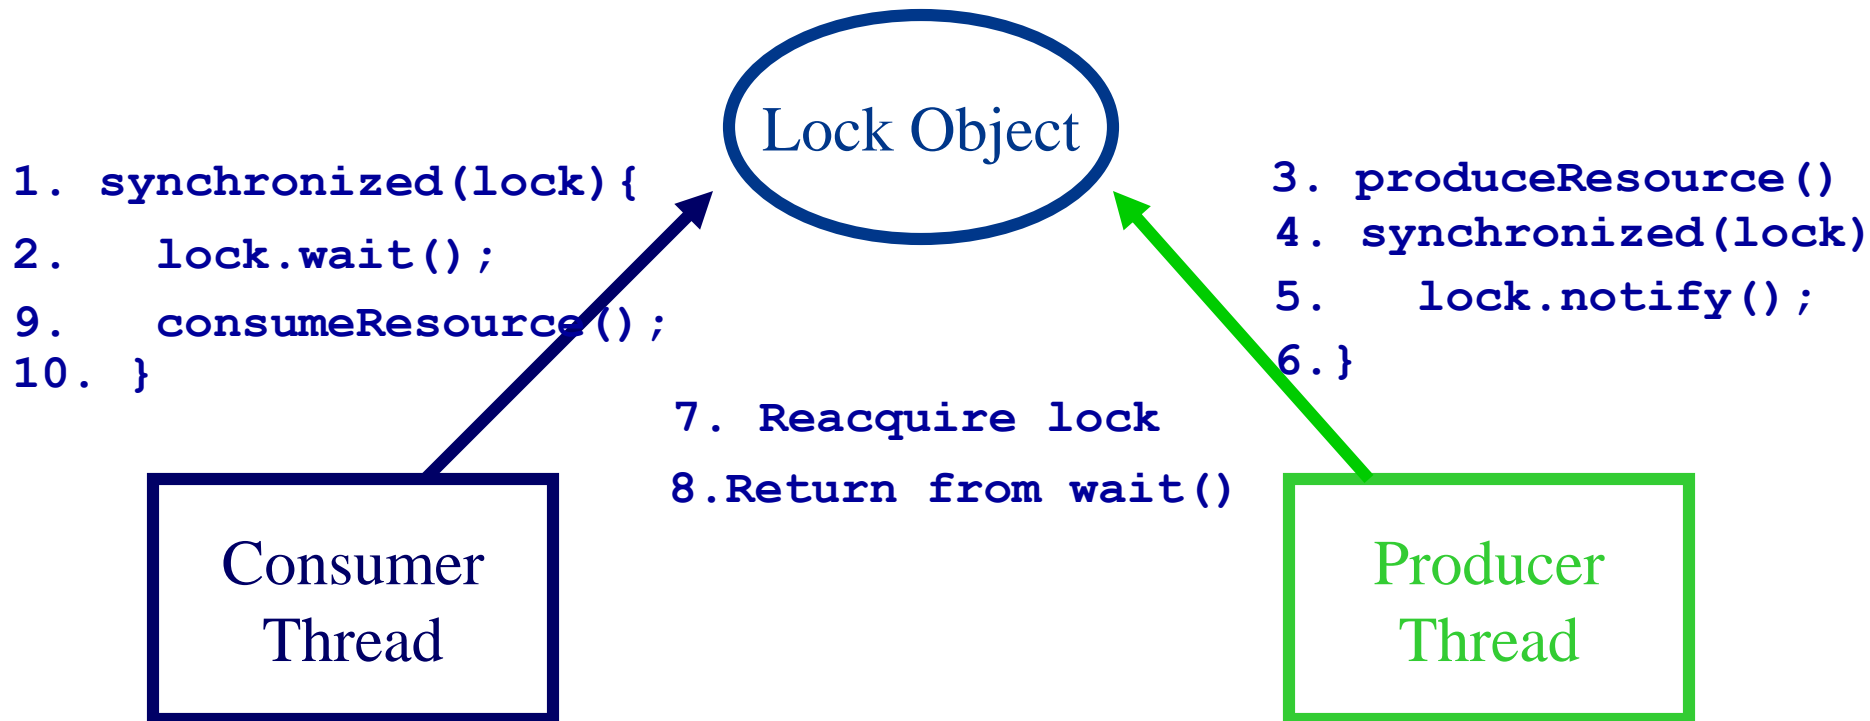
```
synchronized (lock) {  
    while (!resourceAvailable()) {  
        lock.wait();  
    }  
    consumeResource();  
}
```

Producer

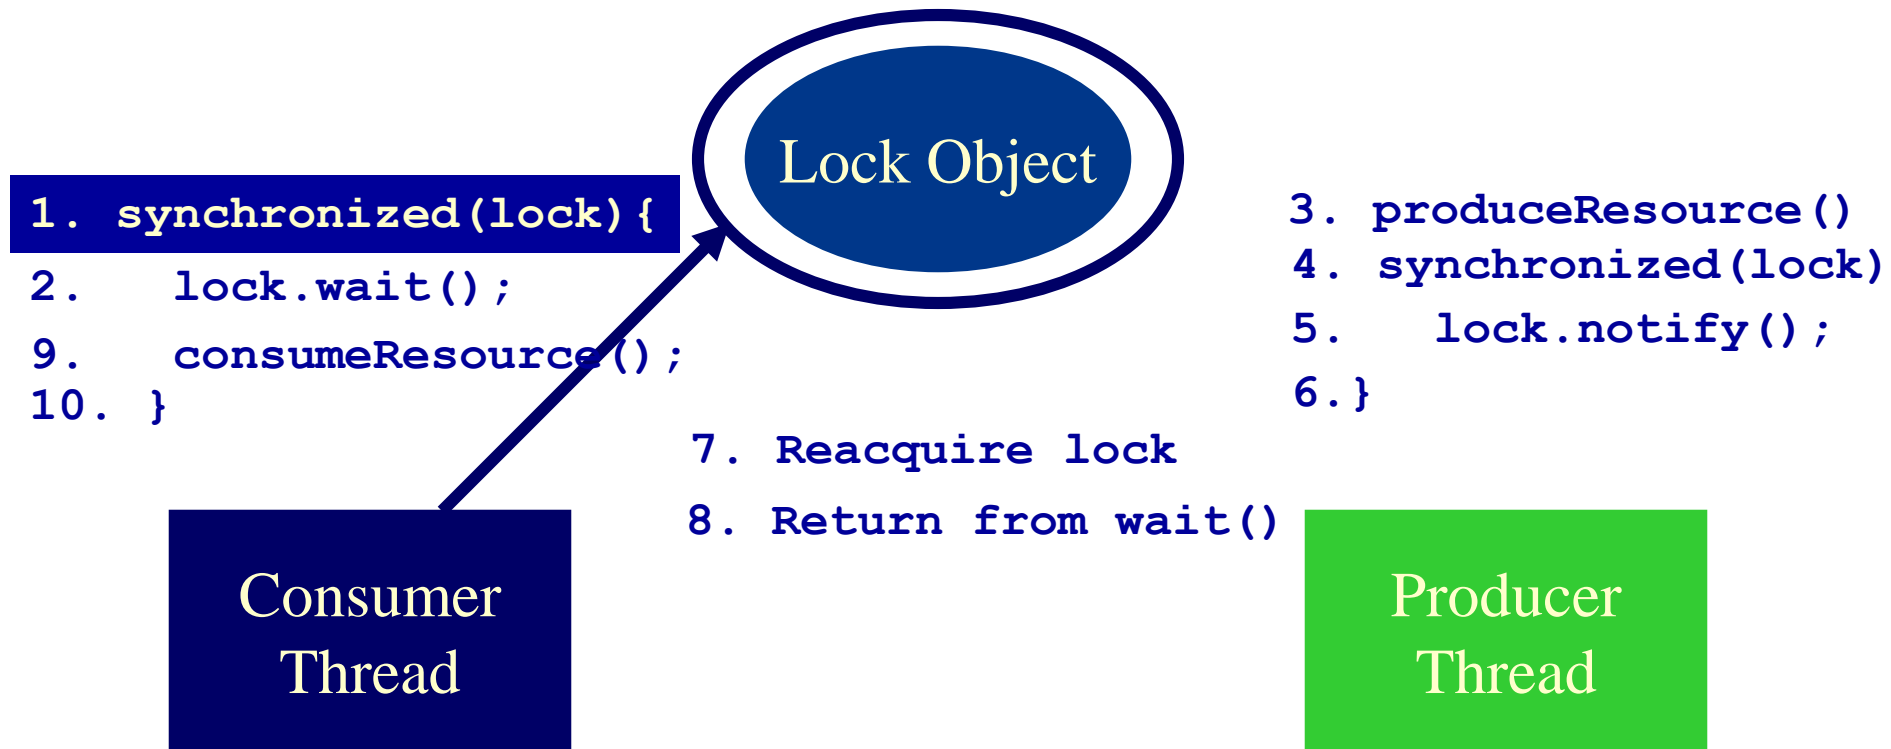
- Producer:

```
produceResource() ;  
synchronized (lock) {  
    lock.notifyAll() ;  
}
```

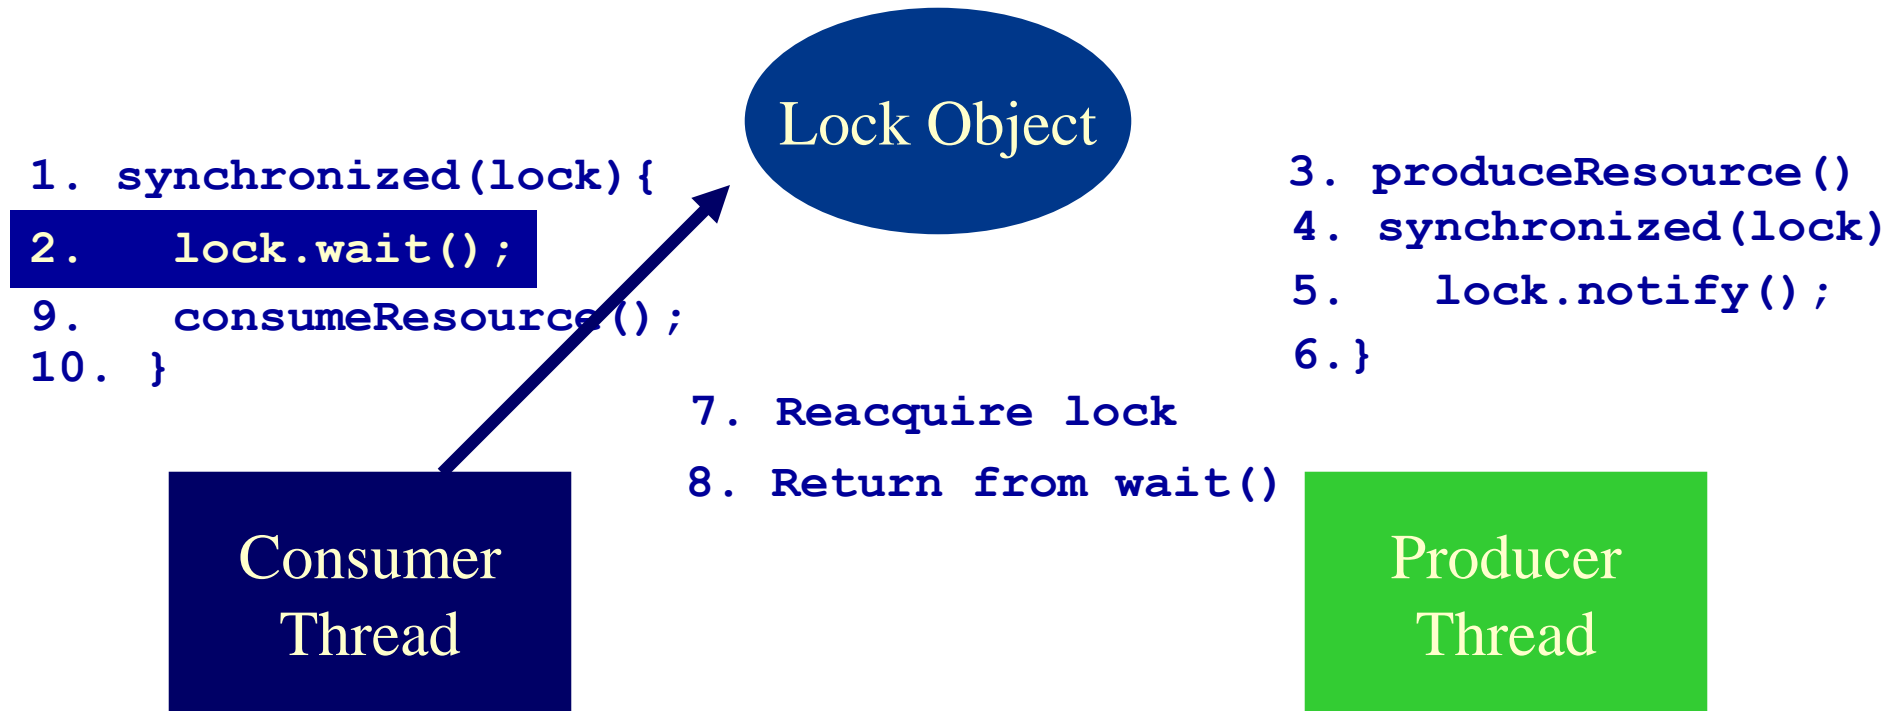
Wait/Notify Sequence



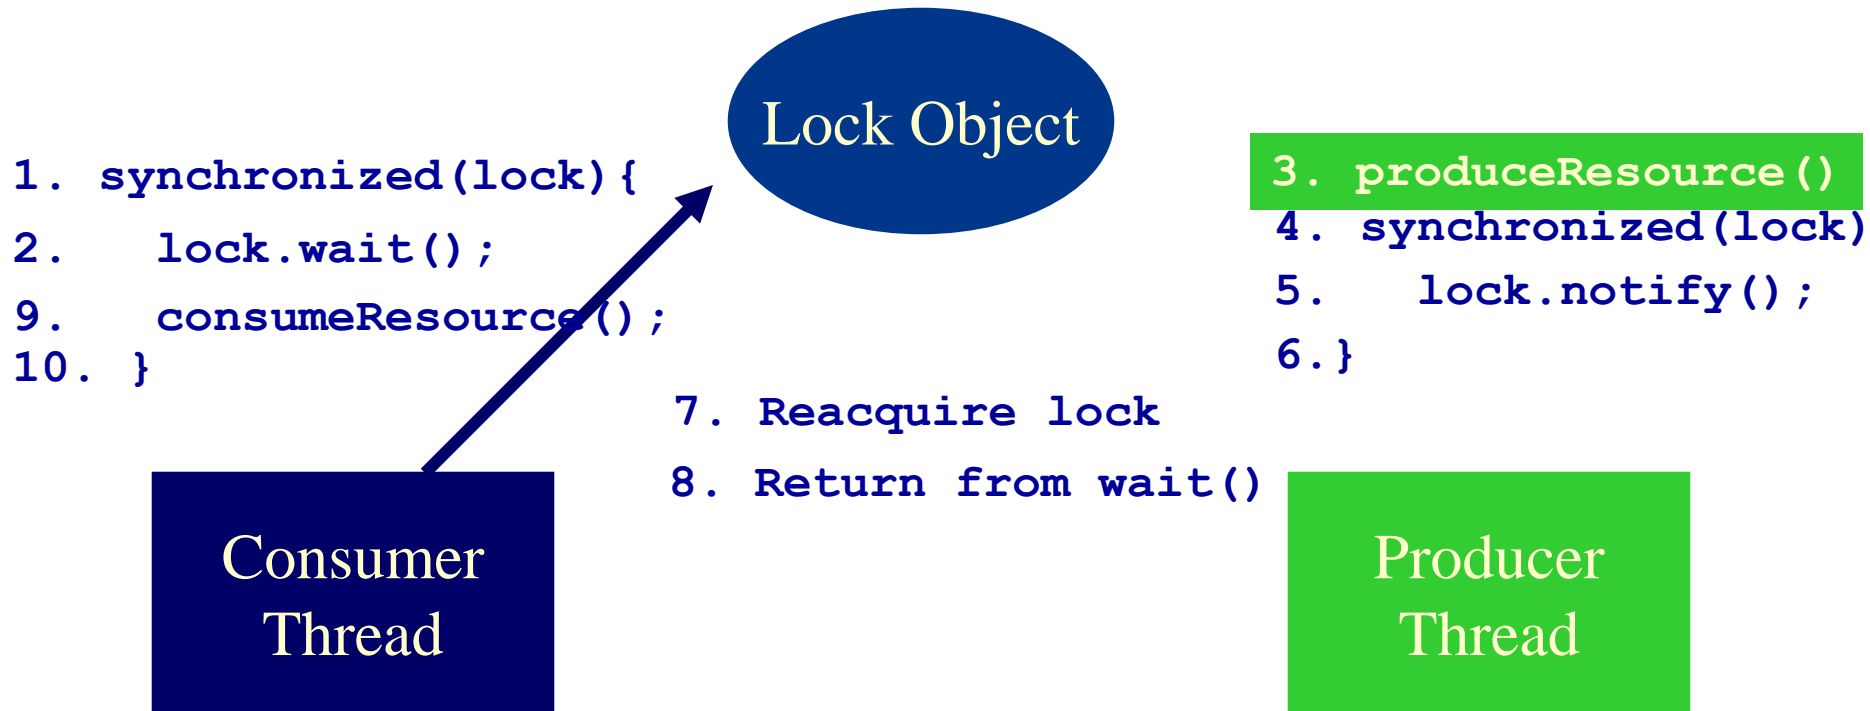
Wait/Notify Sequence



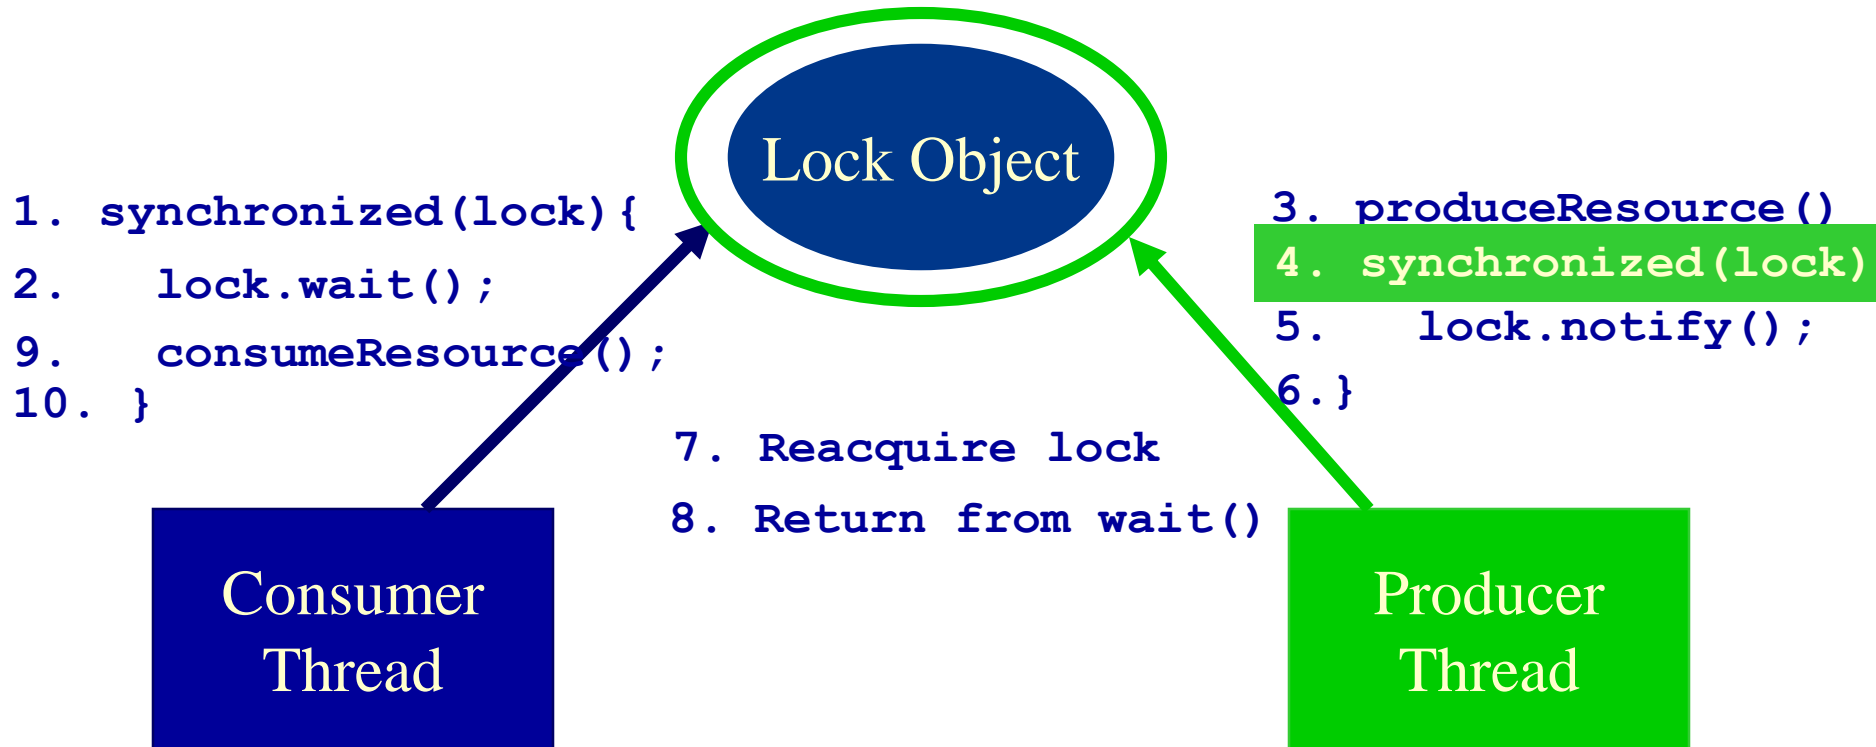
Wait/Notify Sequence



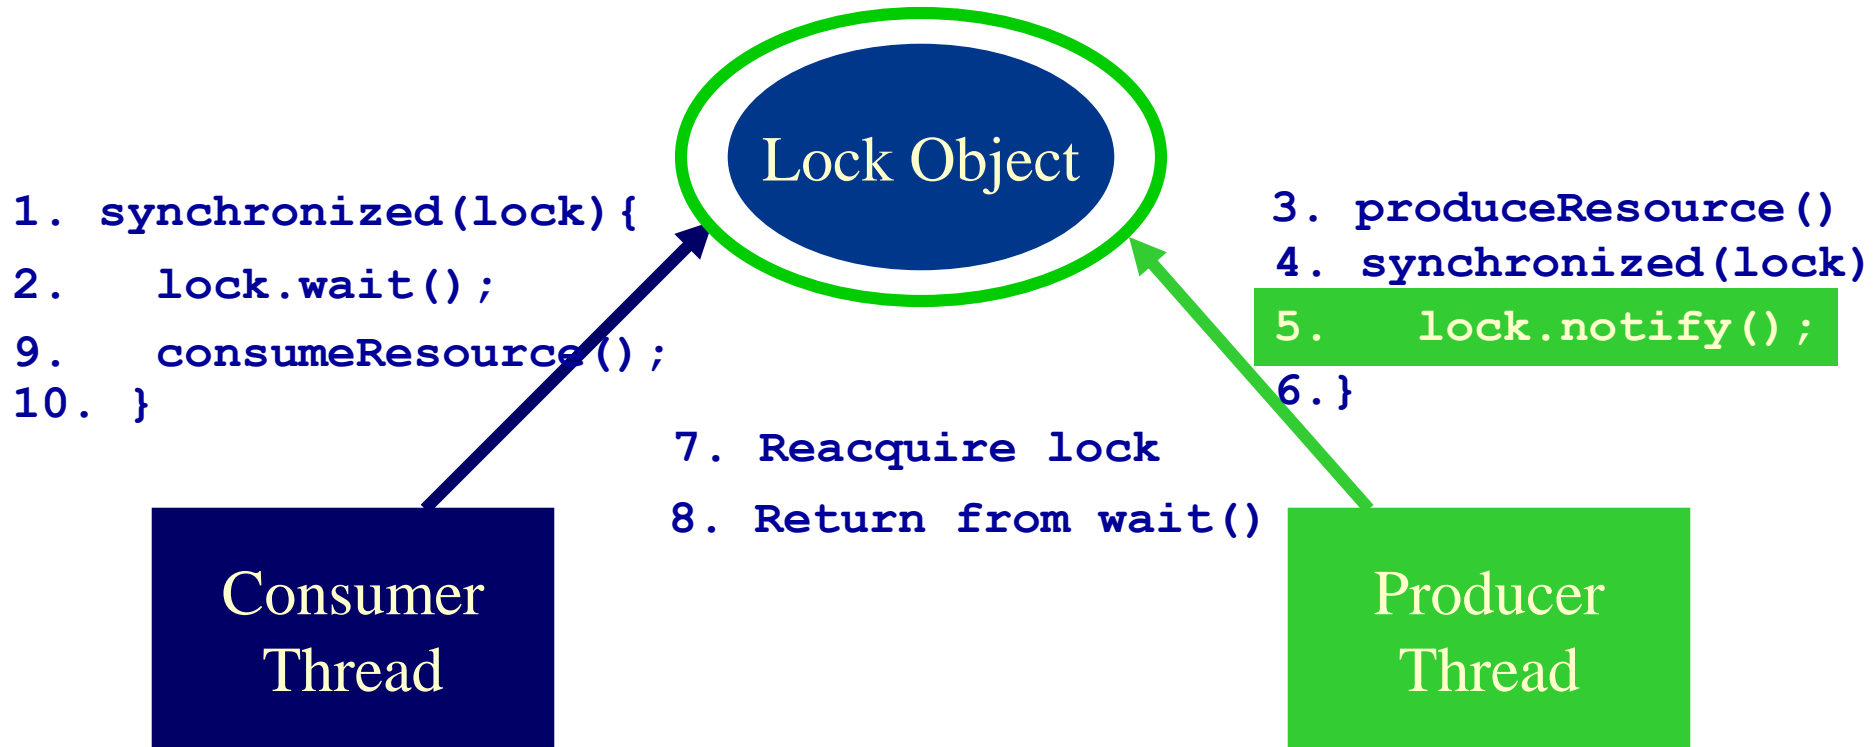
Wait/Notify Sequence



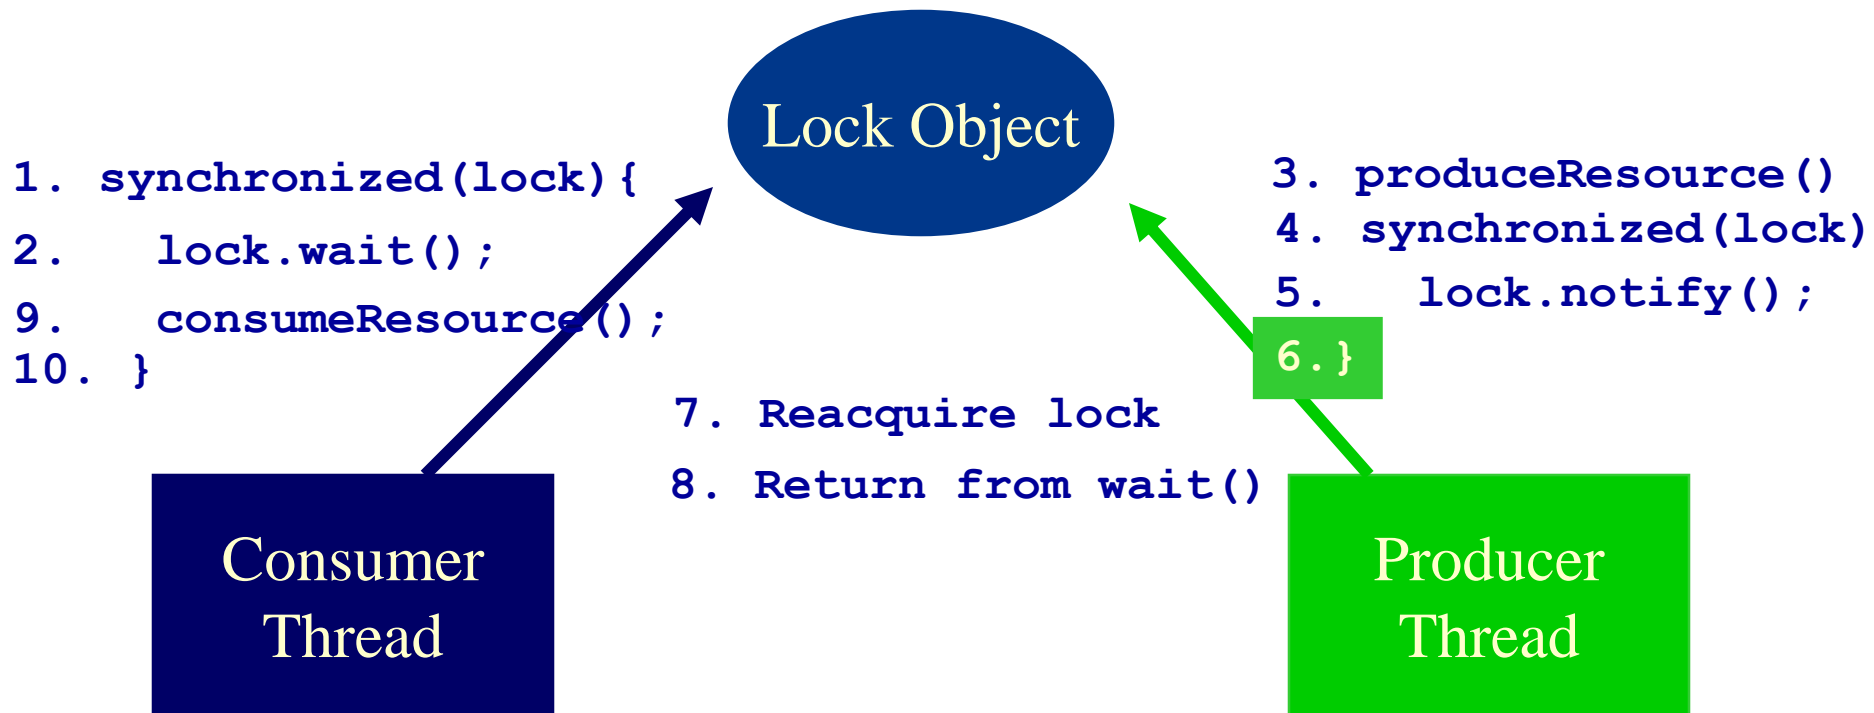
Wait/Notify Sequence




Wait/Notify Sequence




Wait/Notify Sequence



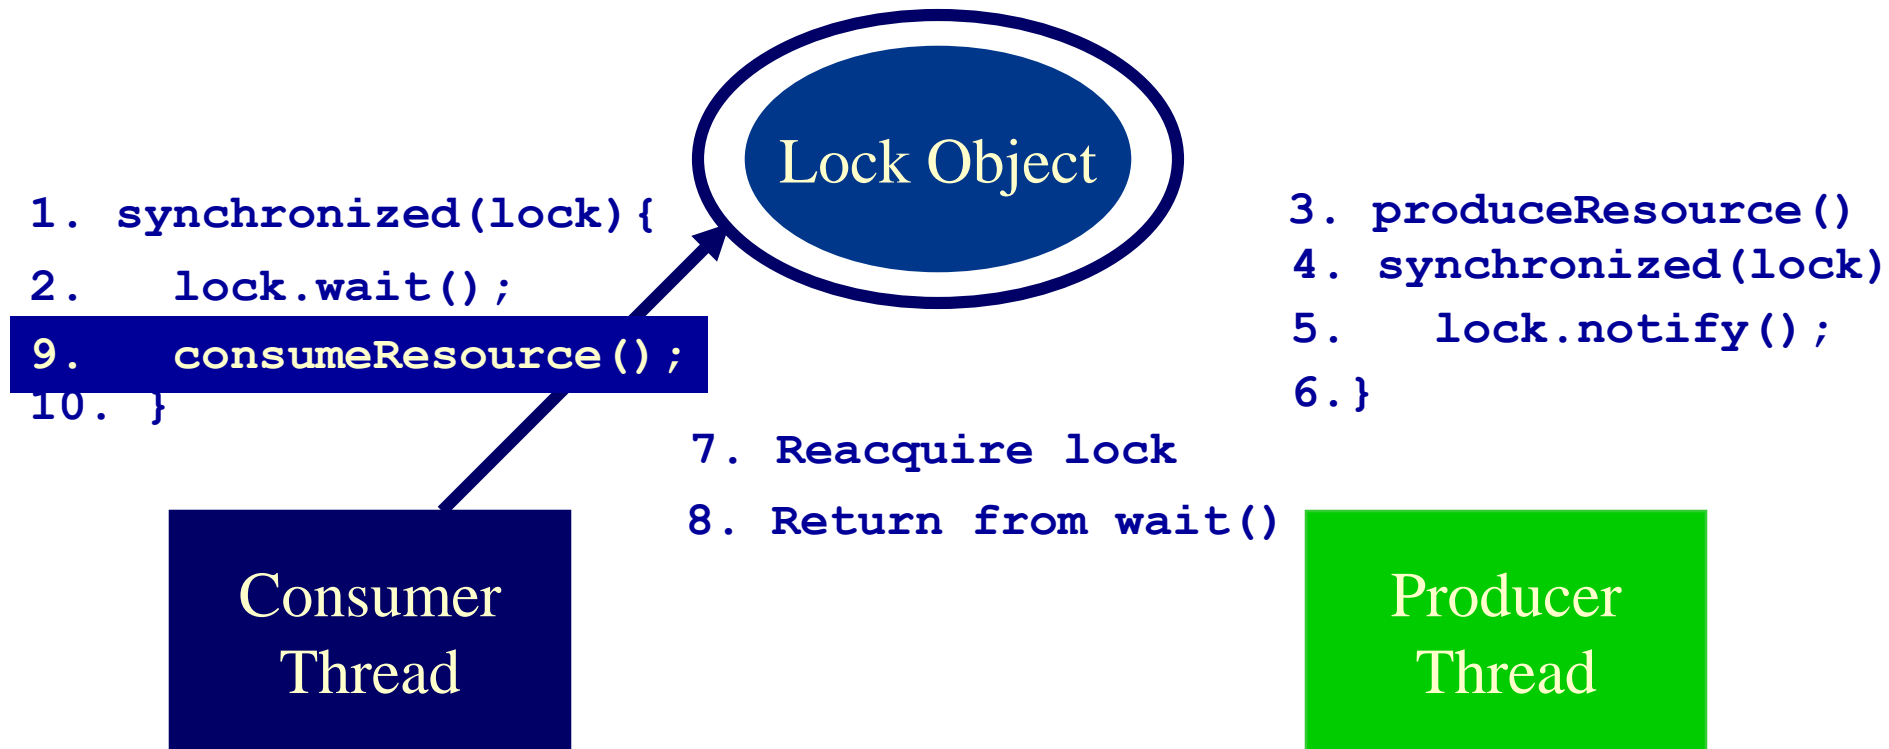
Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence

Lock Object

```
1. synchronized(lock) {  
2.     lock.wait();  
9.     consumeResource();  
10. }
```

Consumer
Thread

```
7. Reacquire lock  
8. Return from wait()
```

```
3. produceResource()  
4. synchronized(lock)  
5.     lock.notify();  
6. }
```

Producer
Thread

The Simpsons Scenario: SimpsonsTest

```
public class SimpsonsTest {  
    public static void main(String[] args) {  
        CookieJar jar = new CookieJar();  
        Homer homer = new Homer(jar);  
        Marge marge = new Marge(jar);  
        new Thread(homer).start();  
        new Thread(marge).start();  
    }  
}
```


The Simpsons Scenario: Homer

```
public class Homer implements Runnable {  
    CookyJar jar;  
  
    public Homer(CookyJar jar) {  
        this.jar = jar;  
    }  
  
    public void eat() {  
        jar.getCooky("Homer");  
        try {  
            Thread.sleep((int)Math.random() * 1000);  
        } catch (InterruptedException ie) {}  
    }  
  
    public void run() {  
        for (int i = 1 ; i <= 10 ; i++) eat();  
    }  
}
```

The Simpsons Scenario: Marge

```
public class Marge implements Runnable {  
    CookieJar jar;  
  
    public Marge(CookieJar jar) {  
        this.jar = jar;  
    }  
  
    public void bake(int cookyNumber) {  
        jar.putCooky("Marge", cookyNumber);  
        try {  
            Thread.sleep((int)Math.random() * 500);  
        } catch (InterruptedException ie) {}  
    }  
  
    public void run() {  
        for (int i = 0 ; i < 10 ; i++) bake(i);  
    }  
}
```

The Simpsons Scenario: CookieJar

```
public class CookieJar {  
    private int contents;  
    private boolean available = false;  
  
    public synchronized void getCookie(String who) {  
        while (!available) {  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        }  
        available = false;  
        notifyAll();  
        System.out.println( who + " ate cookie " +  
                             contents);  
    }  
}
```

The Simpsons Scenario: CookieJar

```
public synchronized void putCooky(String who,
                                   int value) {
    while (available) {
        try {
            wait();
        } catch (InterruptedException e) { }
    }
    contents = value;
    available = true;
    System.out.println(who + " put cooky " +
                       contents + " in the jar");
    notifyAll();
}
}
```

Timers and TimerTask

- The classes `Timer` and `TimerTask` are part of the `java.util` package
- Useful for
 - performing a task after a specified delay
 - performing a sequence of tasks at constant time intervals

Scheduling Timers

- The schedule method of a timer can get as parameters:
 - Task, time
 - Task, time, period
 - Task, delay
 - Task, delay, period

What to do	When to start	At which rate
------------	---------------	---------------

Timer Example

```
import java.util.*;

public class CoffeeTask extends TimerTask {

    public void run() {
        System.out.println("Time for a Coffee Break");
    }

    public static void main(String args[]) {
        Timer timer = new Timer();
        long hour = 1000 * 60 * 60;
        timer.schedule(new CoffeeTask(), 0, 8 * hour);
        timer.scheduleAtFixedRate(new CoffeeTask(),
                                   new Date(), 24 * hour);
    }
}
```

Stopping Timers

- A Timer thread can be stopped in the following ways:
 - Apply `cancel()` on the timer
 - Make the thread a daemon
 - Remove all references to the timer after all the `TimerTask` tasks have finished
 - Call `System.exit()`